

# Deep learning project

Department of Engineering - Aarhus University  
June 3, 2017

Stud. no.: 201270782

Torben Werenberg Vogt

---

Stud. no.: 201270097

Simon Østergaard Kristensen

---

Stud. no.: 201270278

Ivan Bjerring Hansen

---

# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Applied theory</b>	<b>3</b>
<b>3</b>	<b>Implementation</b>	<b>4</b>
<b>4</b>	<b>Results</b>	<b>5</b>
<b>5</b>	<b>Discussion</b>	<b>6</b>
<b>6</b>	<b>Conclusion</b>	<b>7</b>
	<b>Bibliography</b>	<b>8</b>
<b>7</b>	<b>Exercise 1</b>	<b>9</b>
<b>8</b>	<b>Exercise 2</b>	<b>71</b>

# 1 Introduction

## 2 Applied theory

## 3 Implementation

## 4 Results

## 5 Discussion

## 6 Conclusion



# Bibliography

## 7 Exercise 1

# Implementing a Neural Network

In this exercise we will develop a neural network with fully-connected layers to perform classification, and test it out on the CIFAR-10 dataset.

```
In [2]: # A bit of setup

import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The neural network parameters will be stored in a dictionary (model below), where the keys are the parameter names and the values are numpy arrays. Below, we initialize toy data and a toy model that we will use to verify your implementations.

```
In [3]: # Create some toy data to check your implementations
input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5

def init_toy_model():
    model = {}
    model['W1'] = np.linspace(-0.2, 0.6, num=input_size*hidden_size).reshape(input_size, hidden_size)
    model['b1'] = np.linspace(-0.3, 0.7, num=hidden_size)
    model['W2'] = np.linspace(-0.4, 0.1, num=hidden_size*num_classes).reshape(hidden_size, num_classes)
    model['b2'] = np.linspace(-0.5, 0.9, num=num_classes)
    return model

def init_toy_data():
    X = np.linspace(-0.2, 0.5, num=num_inputs*input_size).reshape(num_inputs, input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

model = init_toy_model()
X, y = init_toy_data()
```

## Forward pass: compute scores

Open the file `cs231n/classifiers/neural_net.py` and look at the function `two_layer_net`. This function is very similar to the loss functions you have written for the SVM and Softmax exercises: It takes the data and weights and computes the class scores, the loss, and the gradients on the parameters.

Implement the first part of the forward pass which uses the weights and biases to compute the scores for all inputs.

```
In [4]: from cs231n.classifiers.neural_net import two_layer_net

scores = two_layer_net(X, model)
print scores
correct_scores = [[-0.5328368, 0.20031504, 0.93346689],
                  [-0.59412164, 0.15498488, 0.9040914 ],
                  [-0.67658362, 0.08978957, 0.85616275],
                  [-0.77092643, 0.01339997, 0.79772637],
                  [-0.89110401, -0.08754544, 0.71601312]]

# the difference should be very small. We get 3e-8
print 'Difference between your scores and correct scores:'
print np.sum(np.abs(scores - correct_scores))

[[-0.5328368  0.20031504  0.93346689]
 [-0.59412164  0.15498488  0.9040914 ]
 [-0.67658362  0.08978957  0.85616275]
 [-0.77092643  0.01339997  0.79772637]
 [-0.89110401 -0.08754544  0.71601312]]
Difference between your scores and correct scores:
3.84868227808e-08
```

## Forward pass: compute loss

In the same function, implement the second part that computes the data and regularization loss.

```
In [5]: reg = 0.1
loss, _ = two_layer_net(X, model, y, reg)
correct_loss = 1.38191946092

# should be very small, we get 5e-12
print 'Difference between your loss and correct loss:'
print np.sum(np.abs(loss - correct_loss))

Difference between your loss and correct loss:
4.67736960275e-12
```

## Backward pass

Implement the rest of the function. This will compute the gradient of the loss with respect to the variables  $w_1$ ,  $b_1$ ,  $w_2$ , and  $b_2$ . Now that you (hopefully!) have a correctly implemented forward pass, you can debug your backward pass using a numeric gradient check:

```
In [6]: from cs231n.gradient_check import eval_numerical_gradient

# Use numeric gradient checking to check your implementation of the backward pass.
# If your implementation is correct, the difference between the numeric and
# analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.

loss, grads = two_layer_net(X, model, y, reg)

# these should all be less than 1e-8 or so
for param_name in grads:
    param_grad_num = eval_numerical_gradient(lambda W: two_layer_net(X, model,
y, reg)[0], model[param_name], verbose=False)
    print '%s max relative error: %e' % (param_name, rel_error(param_grad_num, grads[param_name]))

W1 max relative error: 4.426512e-09
W2 max relative error: 1.401432e-09
b2 max relative error: 7.311059e-11
b1 max relative error: 2.746125e-08
```

## Train the network

To train the network we will use SGD with Momentum. Last assignment you implemented vanilla SGD. You will now implement the momentum update and the RMSProp update. Open the file `classifier_trainer.py` and familiarize yourself with the `ClassifierTrainer` class. It performs optimization given an arbitrary cost function data, and model. By default it uses vanilla SGD, which we have already implemented for you. First, run the optimization below using Vanilla SGD:

```

In [7]: from cs231n.classifier_trainer import ClassifierTrainer

model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
                                              reg=0.001,
                                              learning_rate=1e-1, momentum=0.0,
                                              learning_rate_decay=1,
                                              update='sgd', sample_batches=False,
                                              num_epochs=100,
                                              verbose=False)
print 'Final loss with vanilla SGD: %f' % (loss_history[-1], )

starting iteration 0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with vanilla SGD: 0.940686

```

Now fill in the **momentum update** in the first missing code block inside the train function, and run the same optimization as above but with the momentum update. You should see a much better result in the final obtained loss:

```

In [8]: model = init_toy_model()
        trainer = ClassifierTrainer()
        # call the trainer to optimize the loss
        # Notice that we're using sample_batches=False, so we're performing Gradient Descent (no sampled batches of data)
        best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                                    model, two_layer_net,
                                                    reg=0.001,
                                                    learning_rate=1e-1, momentum=0.9,
                                                    learning_rate_decay=1,
                                                    update='momentum',
                                                    sample_batches=False,
                                                    num_epochs=100,
                                                    verbose=False)
        correct_loss = 0.494394
        print 'Final loss with momentum SGD: %f. We get: %f' % (loss_history[-1], correct_loss)

starting iteration 0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with momentum SGD: 0.494394. We get: 0.494394

```

Now also implement the **RMSProp** update rule inside the train function and rerun the optimization:



```
In [9]: model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
                                              reg=0.001,
                                              learning_rate=1e-1, momentum=0.9,
                                              learning_rate_decay=1,
                                              update='rmsprop',
                                              sample_batches=False,
                                              num_epochs=100,
                                              verbose=False)

correct_loss = 0.439368
print 'Final loss with RMSProp: %f. We get: %f' % (loss_history[-1], correct_loss)

starting iteration 0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with RMSProp: 1.125201. We get: 0.439368
```

## Load the data

Now that you have implemented a two-layer network that passes gradient checks, it's time to load up our favorite CIFAR-10 data so we can use it to train a classifier.

```
In [10]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier. These are the same steps as
    we used for the SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = range(num_training, num_training + num_validation)
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = range(num_training)
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = range(num_test)
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis=0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image

    # Reshape data to rows
    X_train = X_train.reshape(num_training, -1)
    X_val = X_val.reshape(num_validation, -1)
    X_test = X_test.reshape(num_test, -1)

    return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print 'Train data shape: ', X_train.shape
print 'Train labels shape: ', y_train.shape
print 'Validation data shape: ', X_val.shape
print 'Validation labels shape: ', y_val.shape
print 'Test data shape: ', X_test.shape
print 'Test labels shape: ', y_test.shape

Train data shape: (49000L, 3072L)
Train labels shape: (49000L,)
Validation data shape: (1000L, 3072L)
Validation labels shape: (1000L,)
Test data shape: (1000L, 3072L)
Test labels shape: (1000L,)
```

## Train a network

To train our network we will use SGD with momentum. In addition, we will adjust the learning rate with an exponential learning rate schedule as optimization proceeds; after each epoch, we will reduce the learning rate by multiplying it by a decay rate.

```
In [11]: from cs231n.classifiers.neural_net import init_two_layer_model

model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, loss_history, train_acc, val_acc = trainer.train(X_train, y_train,
X_val, y_val,
model, two_layer_net,
num_epochs=5, reg=1.0,
momentum=0.9, learning_rate_decay
= 0.95,
learning_rate=1e-5, verbose=True)
```

```
starting iteration 0
Finished epoch 0 / 5: cost 2.302593, train: 0.087000, val 0.082000, lr 1.0000
00e-05
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
starting iteration 100
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
starting iteration 160
starting iteration 170
starting iteration 180
starting iteration 190
starting iteration 200
starting iteration 210
starting iteration 220
starting iteration 230
starting iteration 240
starting iteration 250
starting iteration 260
starting iteration 270
starting iteration 280
starting iteration 290
starting iteration 300
starting iteration 310
starting iteration 320
starting iteration 330
starting iteration 340
starting iteration 350
starting iteration 360
starting iteration 370
starting iteration 380
starting iteration 390
starting iteration 400
starting iteration 410
starting iteration 420
starting iteration 430
starting iteration 440
starting iteration 450
starting iteration 460
starting iteration 470
starting iteration 480
Finished epoch 1 / 5: cost 2.268128, train: 0.158000, val 0.163000, lr 9.5000
00e-06
starting iteration 490
starting iteration 500
starting iteration 510
starting iteration 520
```

```
starting iteration 530
starting iteration 540
starting iteration 550
starting iteration 560
starting iteration 570
starting iteration 580
starting iteration 590
starting iteration 600
starting iteration 610
starting iteration 620
starting iteration 630
starting iteration 640
starting iteration 650
starting iteration 660
starting iteration 670
starting iteration 680
starting iteration 690
starting iteration 700
starting iteration 710
starting iteration 720
starting iteration 730
starting iteration 740
starting iteration 750
starting iteration 760
starting iteration 770
starting iteration 780
starting iteration 790
starting iteration 800
starting iteration 810
starting iteration 820
starting iteration 830
starting iteration 840
starting iteration 850
starting iteration 860
starting iteration 870
starting iteration 880
starting iteration 890
starting iteration 900
starting iteration 910
starting iteration 920
starting iteration 930
starting iteration 940
starting iteration 950
starting iteration 960
starting iteration 970
Finished epoch 2 / 5: cost 1.971051, train: 0.232000, val 0.243000, lr 9.0250
00e-06
starting iteration 980
starting iteration 990
starting iteration 1000
starting iteration 1010
starting iteration 1020
starting iteration 1030
starting iteration 1040
starting iteration 1050
starting iteration 1060
starting iteration 1070
```

```
starting iteration 1080
starting iteration 1090
starting iteration 1100
starting iteration 1110
starting iteration 1120
starting iteration 1130
starting iteration 1140
starting iteration 1150
starting iteration 1160
starting iteration 1170
starting iteration 1180
starting iteration 1190
starting iteration 1200
starting iteration 1210
starting iteration 1220
starting iteration 1230
starting iteration 1240
starting iteration 1250
starting iteration 1260
starting iteration 1270
starting iteration 1280
starting iteration 1290
starting iteration 1300
starting iteration 1310
starting iteration 1320
starting iteration 1330
starting iteration 1340
starting iteration 1350
starting iteration 1360
starting iteration 1370
starting iteration 1380
starting iteration 1390
starting iteration 1400
starting iteration 1410
starting iteration 1420
starting iteration 1430
starting iteration 1440
starting iteration 1450
starting iteration 1460
Finished epoch 3 / 5: cost 1.950968, train: 0.283000, val 0.298000, lr 8.5737
50e-06
starting iteration 1470
starting iteration 1480
starting iteration 1490
starting iteration 1500
starting iteration 1510
starting iteration 1520
starting iteration 1530
starting iteration 1540
starting iteration 1550
starting iteration 1560
starting iteration 1570
starting iteration 1580
starting iteration 1590
starting iteration 1600
starting iteration 1610
starting iteration 1620
```

starting iteration 1630  
starting iteration 1640  
starting iteration 1650  
starting iteration 1660  
starting iteration 1670  
starting iteration 1680  
starting iteration 1690  
starting iteration 1700  
starting iteration 1710  
starting iteration 1720  
starting iteration 1730  
starting iteration 1740  
starting iteration 1750  
starting iteration 1760  
starting iteration 1770  
starting iteration 1780  
starting iteration 1790  
starting iteration 1800  
starting iteration 1810  
starting iteration 1820  
starting iteration 1830  
starting iteration 1840  
starting iteration 1850  
starting iteration 1860  
starting iteration 1870  
starting iteration 1880  
starting iteration 1890  
starting iteration 1900  
starting iteration 1910  
starting iteration 1920  
starting iteration 1930  
starting iteration 1940  
starting iteration 1950  
Finished epoch 4 / 5: cost 1.877080, train: 0.337000, val 0.344000, lr 8.1450  
63e-06  
starting iteration 1960  
starting iteration 1970  
starting iteration 1980  
starting iteration 1990  
starting iteration 2000  
starting iteration 2010  
starting iteration 2020  
starting iteration 2030  
starting iteration 2040  
starting iteration 2050  
starting iteration 2060  
starting iteration 2070  
starting iteration 2080  
starting iteration 2090  
starting iteration 2100  
starting iteration 2110  
starting iteration 2120  
starting iteration 2130  
starting iteration 2140  
starting iteration 2150  
starting iteration 2160  
starting iteration 2170



```
starting iteration 2180
starting iteration 2190
starting iteration 2200
starting iteration 2210
starting iteration 2220
starting iteration 2230
starting iteration 2240
starting iteration 2250
starting iteration 2260
starting iteration 2270
starting iteration 2280
starting iteration 2290
starting iteration 2300
starting iteration 2310
starting iteration 2320
starting iteration 2330
starting iteration 2340
starting iteration 2350
starting iteration 2360
starting iteration 2370
starting iteration 2380
starting iteration 2390
starting iteration 2400
starting iteration 2410
starting iteration 2420
starting iteration 2430
starting iteration 2440
Finished epoch 5 / 5: cost 1.957579, train: 0.331000, val 0.376000, lr 7.7378
09e-06
finished optimization. best validation accuracy: 0.376000
```

## Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.37 on the validation set. This isn't very good.


One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
In [12]: # Plot the Loss function and train / validation accuracies
plt.subplot(2, 1, 1)
plt.plot(loss_history)
plt.title('Loss history')
plt.xlabel('Iteration')
plt.ylabel('Loss')

plt.subplot(2, 1, 2)
plt.plot(train_acc)
plt.plot(val_acc)
plt.legend(['Training accuracy', 'Validation accuracy'], loc='lower right')
plt.xlabel('Epoch')
plt.ylabel('Clasification accuracy')
```

Out[12]: <matplotlib.text.Text at 0x8053ef0>




```
In [13]: from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(model):
    plt.imshow(visualize_grid(model['W1'].T.reshape(-1, 32, 32, 3),
padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(model)
```



# Tune your hyperparameters

**What's wrong?** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the momentum and learning rate decay parameters, but you should be able to get good performance using the default values.

**Approximate results.** You should be aim to achieve a classification accuracy of greater than 50% on the validation set. Our best network gets over 56% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on CIFAR-10 as you can, with a fully-connected Neural Network. For every 1% above 56% on the Test set we will award you with one extra bonus point. Feel free to implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

[illegible]

```
starting iteration 0
Finished epoch 0 / 28: cost 2.302594, train: 0.132000, val 0.109000, lr 1.500
000e-04
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
starting iteration 100
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
starting iteration 160
starting iteration 170
starting iteration 180
starting iteration 190
starting iteration 200
starting iteration 210
starting iteration 220
starting iteration 230
starting iteration 240
starting iteration 250
starting iteration 260
starting iteration 270
starting iteration 280
starting iteration 290
starting iteration 300
starting iteration 310
starting iteration 320
starting iteration 330
starting iteration 340
starting iteration 350
starting iteration 360
starting iteration 370
starting iteration 380
starting iteration 390
starting iteration 400
starting iteration 410
starting iteration 420
starting iteration 430
starting iteration 440
starting iteration 450
starting iteration 460
starting iteration 470
starting iteration 480
Finished epoch 1 / 28: cost 1.677143, train: 0.432000, val 0.416000, lr 1.350
000e-04
starting iteration 490
starting iteration 500
starting iteration 510
starting iteration 520
```

starting iteration 530  
starting iteration 540  
starting iteration 550  
starting iteration 560  
starting iteration 570  
starting iteration 580  
starting iteration 590  
starting iteration 600  
starting iteration 610  
starting iteration 620  
starting iteration 630  
starting iteration 640  
starting iteration 650  
starting iteration 660  
starting iteration 670  
starting iteration 680  
starting iteration 690  
starting iteration 700  
starting iteration 710  
starting iteration 720  
starting iteration 730  
starting iteration 740  
starting iteration 750  
starting iteration 760  
starting iteration 770  
starting iteration 780  
starting iteration 790  
starting iteration 800  
starting iteration 810  
starting iteration 820  
starting iteration 830  
starting iteration 840  
starting iteration 850  
starting iteration 860  
starting iteration 870  
starting iteration 880  
starting iteration 890  
starting iteration 900  
starting iteration 910  
starting iteration 920  
starting iteration 930  
starting iteration 940  
starting iteration 950  
starting iteration 960  
starting iteration 970  
Finished epoch 2 / 28: cost 1.483524, train: 0.496000, val 0.454000, lr 1.215  
000e-04  
starting iteration 980  
starting iteration 990  
starting iteration 1000  
starting iteration 1010  
starting iteration 1020  
starting iteration 1030  
starting iteration 1040  
starting iteration 1050  
starting iteration 1060  
starting iteration 1070

starting iteration 1080  
starting iteration 1090  
starting iteration 1100  
starting iteration 1110  
starting iteration 1120  
starting iteration 1130  
starting iteration 1140  
starting iteration 1150  
starting iteration 1160  
starting iteration 1170  
starting iteration 1180  
starting iteration 1190  
starting iteration 1200  
starting iteration 1210  
starting iteration 1220  
starting iteration 1230  
starting iteration 1240  
starting iteration 1250  
starting iteration 1260  
starting iteration 1270  
starting iteration 1280  
starting iteration 1290  
starting iteration 1300  
starting iteration 1310  
starting iteration 1320  
starting iteration 1330  
starting iteration 1340  
starting iteration 1350  
starting iteration 1360  
starting iteration 1370  
starting iteration 1380  
starting iteration 1390  
starting iteration 1400  
starting iteration 1410  
starting iteration 1420  
starting iteration 1430  
starting iteration 1440  
starting iteration 1450  
starting iteration 1460  
Finished epoch 3 / 28: cost 1.701114, train: 0.433000, val 0.462000, lr 1.093  
500e-04  
starting iteration 1470  
starting iteration 1480  
starting iteration 1490  
starting iteration 1500  
starting iteration 1510  
starting iteration 1520  
starting iteration 1530  
starting iteration 1540  
starting iteration 1550  
starting iteration 1560  
starting iteration 1570  
starting iteration 1580  
starting iteration 1590  
starting iteration 1600  
starting iteration 1610  
starting iteration 1620



starting iteration 1630  
starting iteration 1640  
starting iteration 1650  
starting iteration 1660  
starting iteration 1670  
starting iteration 1680  
starting iteration 1690  
starting iteration 1700  
starting iteration 1710  
starting iteration 1720  
starting iteration 1730  
starting iteration 1740  
starting iteration 1750  
starting iteration 1760  
starting iteration 1770  
starting iteration 1780  
starting iteration 1790  
starting iteration 1800  
starting iteration 1810  
starting iteration 1820  
starting iteration 1830  
starting iteration 1840  
starting iteration 1850  
starting iteration 1860  
starting iteration 1870  
starting iteration 1880  
starting iteration 1890  
starting iteration 1900  
starting iteration 1910  
starting iteration 1920  
starting iteration 1930  
starting iteration 1940  
starting iteration 1950  
Finished epoch 4 / 28: cost 1.704135, train: 0.492000, val 0.454000, lr 9.841500e-05  
starting iteration 1960  
starting iteration 1970  
starting iteration 1980  
starting iteration 1990  
starting iteration 2000  
starting iteration 2010  
starting iteration 2020  
starting iteration 2030  
starting iteration 2040  
starting iteration 2050  
starting iteration 2060  
starting iteration 2070  
starting iteration 2080  
starting iteration 2090  
starting iteration 2100  
starting iteration 2110  
starting iteration 2120  
starting iteration 2130  
starting iteration 2140  
starting iteration 2150  
starting iteration 2160  
starting iteration 2170

starting iteration 2180  
starting iteration 2190  
starting iteration 2200  
starting iteration 2210  
starting iteration 2220  
starting iteration 2230  
starting iteration 2240  
starting iteration 2250  
starting iteration 2260  
starting iteration 2270  
starting iteration 2280  
starting iteration 2290  
starting iteration 2300  
starting iteration 2310  
starting iteration 2320  
starting iteration 2330  
starting iteration 2340  
starting iteration 2350  
starting iteration 2360  
starting iteration 2370  
starting iteration 2380  
starting iteration 2390  
starting iteration 2400  
starting iteration 2410  
starting iteration 2420  
starting iteration 2430  
starting iteration 2440  
Finished epoch 5 / 28: cost 1.551213, train: 0.504000, val 0.486000, lr 8.857  
350e-05  
starting iteration 2450  
starting iteration 2460  
starting iteration 2470  
starting iteration 2480  
starting iteration 2490  
starting iteration 2500  
starting iteration 2510  
starting iteration 2520  
starting iteration 2530  
starting iteration 2540  
starting iteration 2550  
starting iteration 2560  
starting iteration 2570  
starting iteration 2580  
starting iteration 2590  
starting iteration 2600  
starting iteration 2610  
starting iteration 2620  
starting iteration 2630  
starting iteration 2640  
starting iteration 2650  
starting iteration 2660  
starting iteration 2670  
starting iteration 2680  
starting iteration 2690  
starting iteration 2700  
starting iteration 2710  
starting iteration 2720

starting iteration 2730  
starting iteration 2740  
starting iteration 2750  
starting iteration 2760  
starting iteration 2770  
starting iteration 2780  
starting iteration 2790  
starting iteration 2800  
starting iteration 2810  
starting iteration 2820  
starting iteration 2830  
starting iteration 2840  
starting iteration 2850  
starting iteration 2860  
starting iteration 2870  
starting iteration 2880  
starting iteration 2890  
starting iteration 2900  
starting iteration 2910  
starting iteration 2920  
starting iteration 2930  
Finished epoch 6 / 28: cost 1.593545, train: 0.485000, val 0.486000, lr 7.971  
615e-05  
starting iteration 2940  
starting iteration 2950  
starting iteration 2960  
starting iteration 2970  
starting iteration 2980  
starting iteration 2990  
starting iteration 3000  
starting iteration 3010  
starting iteration 3020  
starting iteration 3030  
starting iteration 3040  
starting iteration 3050  
starting iteration 3060  
starting iteration 3070  
starting iteration 3080  
starting iteration 3090  
starting iteration 3100  
starting iteration 3110  
starting iteration 3120  
starting iteration 3130  
starting iteration 3140  
starting iteration 3150  
starting iteration 3160  
starting iteration 3170  
starting iteration 3180  
starting iteration 3190  
starting iteration 3200  
starting iteration 3210  
starting iteration 3220  
starting iteration 3230  
starting iteration 3240  
starting iteration 3250  
starting iteration 3260  
starting iteration 3270

starting iteration 3280  
starting iteration 3290  
starting iteration 3300  
starting iteration 3310  
starting iteration 3320  
starting iteration 3330  
starting iteration 3340  
starting iteration 3350  
starting iteration 3360  
starting iteration 3370  
starting iteration 3380  
starting iteration 3390  
starting iteration 3400  
starting iteration 3410  
starting iteration 3420  
Finished epoch 7 / 28: cost 1.526852, train: 0.502000, val 0.496000, lr 7.174453e-05  
starting iteration 3430  
starting iteration 3440  
starting iteration 3450  
starting iteration 3460  
starting iteration 3470  
starting iteration 3480  
starting iteration 3490  
starting iteration 3500  
starting iteration 3510  
starting iteration 3520  
starting iteration 3530  
starting iteration 3540  
starting iteration 3550  
starting iteration 3560  
starting iteration 3570  
starting iteration 3580  
starting iteration 3590  
starting iteration 3600  
starting iteration 3610  
starting iteration 3620  
starting iteration 3630  
starting iteration 3640  
starting iteration 3650  
starting iteration 3660  
starting iteration 3670  
starting iteration 3680  
starting iteration 3690  
starting iteration 3700  
starting iteration 3710  
starting iteration 3720  
starting iteration 3730  
starting iteration 3740  
starting iteration 3750  
starting iteration 3760  
starting iteration 3770  
starting iteration 3780  
starting iteration 3790  
starting iteration 3800  
starting iteration 3810  
starting iteration 3820

```
starting iteration 3830
starting iteration 3840
starting iteration 3850
starting iteration 3860
starting iteration 3870
starting iteration 3880
starting iteration 3890
starting iteration 3900
starting iteration 3910
Finished epoch 8 / 28: cost 1.754050, train: 0.488000, val 0.505000, lr 6.457
008e-05
starting iteration 3920
starting iteration 3930
starting iteration 3940
starting iteration 3950
starting iteration 3960
starting iteration 3970
starting iteration 3980
starting iteration 3990
starting iteration 4000
starting iteration 4010
starting iteration 4020
starting iteration 4030
starting iteration 4040
starting iteration 4050
starting iteration 4060
starting iteration 4070
starting iteration 4080
starting iteration 4090
starting iteration 4100
starting iteration 4110
starting iteration 4120
starting iteration 4130
starting iteration 4140
starting iteration 4150
starting iteration 4160
starting iteration 4170
starting iteration 4180
starting iteration 4190
starting iteration 4200
starting iteration 4210
starting iteration 4220
starting iteration 4230
starting iteration 4240
starting iteration 4250
starting iteration 4260
starting iteration 4270
starting iteration 4280
starting iteration 4290
starting iteration 4300
starting iteration 4310
starting iteration 4320
starting iteration 4330
starting iteration 4340
starting iteration 4350
starting iteration 4360
starting iteration 4370
```

```
starting iteration 4380
starting iteration 4390
starting iteration 4400
Finished epoch 9 / 28: cost 1.611458, train: 0.513000, val 0.486000, lr 5.811
307e-05
starting iteration 4410
starting iteration 4420
starting iteration 4430
starting iteration 4440
starting iteration 4450
starting iteration 4460
starting iteration 4470
starting iteration 4480
starting iteration 4490
starting iteration 4500
starting iteration 4510
starting iteration 4520
starting iteration 4530
starting iteration 4540
starting iteration 4550
starting iteration 4560
starting iteration 4570
starting iteration 4580
starting iteration 4590
starting iteration 4600
starting iteration 4610
starting iteration 4620
starting iteration 4630
starting iteration 4640
starting iteration 4650
starting iteration 4660
starting iteration 4670
starting iteration 4680
starting iteration 4690
starting iteration 4700
starting iteration 4710
starting iteration 4720
starting iteration 4730
starting iteration 4740
starting iteration 4750
starting iteration 4760
starting iteration 4770
starting iteration 4780
starting iteration 4790
starting iteration 4800
starting iteration 4810
starting iteration 4820
starting iteration 4830
starting iteration 4840
starting iteration 4850
starting iteration 4860
starting iteration 4870
starting iteration 4880
starting iteration 4890
Finished epoch 10 / 28: cost 1.579166, train: 0.502000, val 0.499000, lr 5.23
0177e-05
starting iteration 4900
```

starting iteration 4910  
starting iteration 4920  
starting iteration 4930  
starting iteration 4940  
starting iteration 4950  
starting iteration 4960  
starting iteration 4970  
starting iteration 4980  
starting iteration 4990  
starting iteration 5000  
starting iteration 5010  
starting iteration 5020  
starting iteration 5030  
starting iteration 5040  
starting iteration 5050  
starting iteration 5060  
starting iteration 5070  
starting iteration 5080  
starting iteration 5090  
starting iteration 5100  
starting iteration 5110  
starting iteration 5120  
starting iteration 5130  
starting iteration 5140  
starting iteration 5150  
starting iteration 5160  
starting iteration 5170  
starting iteration 5180  
starting iteration 5190  
starting iteration 5200  
starting iteration 5210  
starting iteration 5220  
starting iteration 5230  
starting iteration 5240  
starting iteration 5250  
starting iteration 5260  
starting iteration 5270  
starting iteration 5280  
starting iteration 5290  
starting iteration 5300  
starting iteration 5310  
starting iteration 5320  
starting iteration 5330  
starting iteration 5340  
starting iteration 5350  
starting iteration 5360  
starting iteration 5370  
starting iteration 5380  
Finished epoch 11 / 28: cost 1.721801, train: 0.530000, val 0.478000, lr 4.70  
7159e-05  
starting iteration 5390  
starting iteration 5400  
starting iteration 5410  
starting iteration 5420  
starting iteration 5430  
starting iteration 5440  
starting iteration 5450

```
starting iteration 5460
starting iteration 5470
starting iteration 5480
starting iteration 5490
starting iteration 5500
starting iteration 5510
starting iteration 5520
starting iteration 5530
starting iteration 5540
starting iteration 5550
starting iteration 5560
starting iteration 5570
starting iteration 5580
starting iteration 5590
starting iteration 5600
starting iteration 5610
starting iteration 5620
starting iteration 5630
starting iteration 5640
starting iteration 5650
starting iteration 5660
starting iteration 5670
starting iteration 5680
starting iteration 5690
starting iteration 5700
starting iteration 5710
starting iteration 5720
starting iteration 5730
starting iteration 5740
starting iteration 5750
starting iteration 5760
starting iteration 5770
starting iteration 5780
starting iteration 5790
starting iteration 5800
starting iteration 5810
starting iteration 5820
starting iteration 5830
starting iteration 5840
starting iteration 5850
starting iteration 5860
starting iteration 5870
Finished epoch 12 / 28: cost 1.628940, train: 0.536000, val 0.496000, lr 4.23
6443e-05
starting iteration 5880
starting iteration 5890
starting iteration 5900
starting iteration 5910
starting iteration 5920
starting iteration 5930
starting iteration 5940
starting iteration 5950
starting iteration 5960
starting iteration 5970
starting iteration 5980
starting iteration 5990
starting iteration 6000
```



```
starting iteration 6010
starting iteration 6020
starting iteration 6030
starting iteration 6040
starting iteration 6050
starting iteration 6060
starting iteration 6070
starting iteration 6080
starting iteration 6090
starting iteration 6100
starting iteration 6110
starting iteration 6120
starting iteration 6130
starting iteration 6140
starting iteration 6150
starting iteration 6160
starting iteration 6170
starting iteration 6180
starting iteration 6190
starting iteration 6200
starting iteration 6210
starting iteration 6220
starting iteration 6230
starting iteration 6240
starting iteration 6250
starting iteration 6260
starting iteration 6270
starting iteration 6280
starting iteration 6290
starting iteration 6300
starting iteration 6310
starting iteration 6320
starting iteration 6330
starting iteration 6340
starting iteration 6350
starting iteration 6360
Finished epoch 13 / 28: cost 1.558282, train: 0.513000, val 0.511000, lr 3.81
2799e-05
starting iteration 6370
starting iteration 6380
starting iteration 6390
starting iteration 6400
starting iteration 6410
starting iteration 6420
starting iteration 6430
starting iteration 6440
starting iteration 6450
starting iteration 6460
starting iteration 6470
starting iteration 6480
starting iteration 6490
starting iteration 6500
starting iteration 6510
starting iteration 6520
starting iteration 6530
starting iteration 6540
starting iteration 6550
```

```
starting iteration 6560
starting iteration 6570
starting iteration 6580
starting iteration 6590
starting iteration 6600
starting iteration 6610
starting iteration 6620
starting iteration 6630
starting iteration 6640
starting iteration 6650
starting iteration 6660
starting iteration 6670
starting iteration 6680
starting iteration 6690
starting iteration 6700
starting iteration 6710
starting iteration 6720
starting iteration 6730
starting iteration 6740
starting iteration 6750
starting iteration 6760
starting iteration 6770
starting iteration 6780
starting iteration 6790
starting iteration 6800
starting iteration 6810
starting iteration 6820
starting iteration 6830
starting iteration 6840
starting iteration 6850
Finished epoch 14 / 28: cost 1.523632, train: 0.546000, val 0.502000, lr 3.43
1519e-05
starting iteration 6860
starting iteration 6870
starting iteration 6880
starting iteration 6890
starting iteration 6900
starting iteration 6910
starting iteration 6920
starting iteration 6930
starting iteration 6940
starting iteration 6950
starting iteration 6960
starting iteration 6970
starting iteration 6980
starting iteration 6990
starting iteration 7000
starting iteration 7010
starting iteration 7020
starting iteration 7030
starting iteration 7040
starting iteration 7050
starting iteration 7060
starting iteration 7070
starting iteration 7080
starting iteration 7090
starting iteration 7100
```

```
starting iteration 7110
starting iteration 7120
starting iteration 7130
starting iteration 7140
starting iteration 7150
starting iteration 7160
starting iteration 7170
starting iteration 7180
starting iteration 7190
starting iteration 7200
starting iteration 7210
starting iteration 7220
starting iteration 7230
starting iteration 7240
starting iteration 7250
starting iteration 7260
starting iteration 7270
starting iteration 7280
starting iteration 7290
starting iteration 7300
starting iteration 7310
starting iteration 7320
starting iteration 7330
starting iteration 7340
Finished epoch 15 / 28: cost 1.396022, train: 0.518000, val 0.489000, lr 3.08
8367e-05
starting iteration 7350
starting iteration 7360
starting iteration 7370
starting iteration 7380
starting iteration 7390
starting iteration 7400
starting iteration 7410
starting iteration 7420
starting iteration 7430
starting iteration 7440
starting iteration 7450
starting iteration 7460
starting iteration 7470
starting iteration 7480
starting iteration 7490
starting iteration 7500
starting iteration 7510
starting iteration 7520
starting iteration 7530
starting iteration 7540
starting iteration 7550
starting iteration 7560
starting iteration 7570
starting iteration 7580
starting iteration 7590
starting iteration 7600
starting iteration 7610
starting iteration 7620
starting iteration 7630
starting iteration 7640
starting iteration 7650
```

starting iteration 7660  
starting iteration 7670  
starting iteration 7680  
starting iteration 7690  
starting iteration 7700  
starting iteration 7710  
starting iteration 7720  
starting iteration 7730  
starting iteration 7740  
starting iteration 7750  
starting iteration 7760  
starting iteration 7770  
starting iteration 7780  
starting iteration 7790  
starting iteration 7800  
starting iteration 7810  
starting iteration 7820  
starting iteration 7830  
Finished epoch 16 / 28: cost 1.521139, train: 0.555000, val 0.509000, lr 2.77  
9530e-05  
starting iteration 7840  
starting iteration 7850  
starting iteration 7860  
starting iteration 7870  
starting iteration 7880  
starting iteration 7890  
starting iteration 7900  
starting iteration 7910  
starting iteration 7920  
starting iteration 7930  
starting iteration 7940  
starting iteration 7950  
starting iteration 7960  
starting iteration 7970  
starting iteration 7980  
starting iteration 7990  
starting iteration 8000  
starting iteration 8010  
starting iteration 8020  
starting iteration 8030  
starting iteration 8040  
starting iteration 8050  
starting iteration 8060  
starting iteration 8070  
starting iteration 8080  
starting iteration 8090  
starting iteration 8100  
starting iteration 8110  
starting iteration 8120  
starting iteration 8130  
starting iteration 8140  
starting iteration 8150  
starting iteration 8160  
starting iteration 8170  
starting iteration 8180  
starting iteration 8190  
starting iteration 8200

```
starting iteration 8210
starting iteration 8220
starting iteration 8230
starting iteration 8240
starting iteration 8250
starting iteration 8260
starting iteration 8270
starting iteration 8280
starting iteration 8290
starting iteration 8300
starting iteration 8310
starting iteration 8320
Finished epoch 17 / 28: cost 1.508517, train: 0.552000, val 0.526000, lr 2.50
1577e-05
starting iteration 8330
starting iteration 8340
starting iteration 8350
starting iteration 8360
starting iteration 8370
starting iteration 8380
starting iteration 8390
starting iteration 8400
starting iteration 8410
starting iteration 8420
starting iteration 8430
starting iteration 8440
starting iteration 8450
starting iteration 8460
starting iteration 8470
starting iteration 8480
starting iteration 8490
starting iteration 8500
starting iteration 8510
starting iteration 8520
starting iteration 8530
starting iteration 8540
starting iteration 8550
starting iteration 8560
starting iteration 8570
starting iteration 8580
starting iteration 8590
starting iteration 8600
starting iteration 8610
starting iteration 8620
starting iteration 8630
starting iteration 8640
starting iteration 8650
starting iteration 8660
starting iteration 8670
starting iteration 8680
starting iteration 8690
starting iteration 8700
starting iteration 8710
starting iteration 8720
starting iteration 8730
starting iteration 8740
starting iteration 8750
```

```
starting iteration 8760
starting iteration 8770
starting iteration 8780
starting iteration 8790
starting iteration 8800
starting iteration 8810
Finished epoch 18 / 28: cost 1.448959, train: 0.536000, val 0.497000, lr 2.25
1420e-05
starting iteration 8820
starting iteration 8830
starting iteration 8840
starting iteration 8850
starting iteration 8860
starting iteration 8870
starting iteration 8880
starting iteration 8890
starting iteration 8900
starting iteration 8910
starting iteration 8920
starting iteration 8930
starting iteration 8940
starting iteration 8950
starting iteration 8960
starting iteration 8970
starting iteration 8980
starting iteration 8990
starting iteration 9000
starting iteration 9010
starting iteration 9020
starting iteration 9030
starting iteration 9040
starting iteration 9050
starting iteration 9060
starting iteration 9070
starting iteration 9080
starting iteration 9090
starting iteration 9100
starting iteration 9110
starting iteration 9120
starting iteration 9130
starting iteration 9140
starting iteration 9150
starting iteration 9160
starting iteration 9170
starting iteration 9180
starting iteration 9190
starting iteration 9200
starting iteration 9210
starting iteration 9220
starting iteration 9230
starting iteration 9240
starting iteration 9250
starting iteration 9260
starting iteration 9270
starting iteration 9280
starting iteration 9290
starting iteration 9300
```

Finished epoch 19 / 28: cost 1.346652, train: 0.557000, val 0.507000, lr 2.026278e-05

starting iteration 9310  
starting iteration 9320  
starting iteration 9330  
starting iteration 9340  
starting iteration 9350  
starting iteration 9360  
starting iteration 9370  
starting iteration 9380  
starting iteration 9390  
starting iteration 9400  
starting iteration 9410  
starting iteration 9420  
starting iteration 9430  
starting iteration 9440  
starting iteration 9450  
starting iteration 9460  
starting iteration 9470  
starting iteration 9480  
starting iteration 9490  
starting iteration 9500  
starting iteration 9510  
starting iteration 9520  
starting iteration 9530  
starting iteration 9540  
starting iteration 9550  
starting iteration 9560  
starting iteration 9570  
starting iteration 9580  
starting iteration 9590  
starting iteration 9600  
starting iteration 9610  
starting iteration 9620  
starting iteration 9630  
starting iteration 9640  
starting iteration 9650  
starting iteration 9660  
starting iteration 9670  
starting iteration 9680  
starting iteration 9690  
starting iteration 9700  
starting iteration 9710  
starting iteration 9720  
starting iteration 9730  
starting iteration 9740  
starting iteration 9750  
starting iteration 9760  
starting iteration 9770  
starting iteration 9780  
starting iteration 9790

Finished epoch 20 / 28: cost 1.515841, train: 0.569000, val 0.506000, lr 1.823650e-05

starting iteration 9800  
starting iteration 9810  
starting iteration 9820  
starting iteration 9830

```
starting iteration 9840
starting iteration 9850
starting iteration 9860
starting iteration 9870
starting iteration 9880
starting iteration 9890
starting iteration 9900
starting iteration 9910
starting iteration 9920
starting iteration 9930
starting iteration 9940
starting iteration 9950
starting iteration 9960
starting iteration 9970
starting iteration 9980
starting iteration 9990
starting iteration 10000
starting iteration 10010
starting iteration 10020
starting iteration 10030
starting iteration 10040
starting iteration 10050
starting iteration 10060
starting iteration 10070
starting iteration 10080
starting iteration 10090
starting iteration 10100
starting iteration 10110
starting iteration 10120
starting iteration 10130
starting iteration 10140
starting iteration 10150
starting iteration 10160
starting iteration 10170
starting iteration 10180
starting iteration 10190
starting iteration 10200
starting iteration 10210
starting iteration 10220
starting iteration 10230
starting iteration 10240
starting iteration 10250
starting iteration 10260
starting iteration 10270
starting iteration 10280
Finished epoch 21 / 28: cost 1.469721, train: 0.549000, val 0.520000, lr 1.64
1285e-05
starting iteration 10290
starting iteration 10300
starting iteration 10310
starting iteration 10320
starting iteration 10330
starting iteration 10340
starting iteration 10350
starting iteration 10360
starting iteration 10370
starting iteration 10380
```



starting iteration 10390  
starting iteration 10400  
starting iteration 10410  
starting iteration 10420  
starting iteration 10430  
starting iteration 10440  
starting iteration 10450  
starting iteration 10460  
starting iteration 10470  
starting iteration 10480  
starting iteration 10490  
starting iteration 10500  
starting iteration 10510  
starting iteration 10520  
starting iteration 10530  
starting iteration 10540  
starting iteration 10550  
starting iteration 10560  
starting iteration 10570  
starting iteration 10580  
starting iteration 10590  
starting iteration 10600  
starting iteration 10610  
starting iteration 10620  
starting iteration 10630  
starting iteration 10640  
starting iteration 10650  
starting iteration 10660  
starting iteration 10670  
starting iteration 10680  
starting iteration 10690  
starting iteration 10700  
starting iteration 10710  
starting iteration 10720  
starting iteration 10730  
starting iteration 10740  
starting iteration 10750  
starting iteration 10760  
starting iteration 10770  
Finished epoch 22 / 28: cost 1.465887, train: 0.552000, val 0.516000, lr 1.47  
7156e-05  
starting iteration 10780  
starting iteration 10790  
starting iteration 10800  
starting iteration 10810  
starting iteration 10820  
starting iteration 10830  
starting iteration 10840  
starting iteration 10850  
starting iteration 10860  
starting iteration 10870  
starting iteration 10880  
starting iteration 10890  
starting iteration 10900  
starting iteration 10910  
starting iteration 10920  
starting iteration 10930

```
starting iteration 10940
starting iteration 10950
starting iteration 10960
starting iteration 10970
starting iteration 10980
starting iteration 10990
starting iteration 11000
starting iteration 11010
starting iteration 11020
starting iteration 11030
starting iteration 11040
starting iteration 11050
starting iteration 11060
starting iteration 11070
starting iteration 11080
starting iteration 11090
starting iteration 11100
starting iteration 11110
starting iteration 11120
starting iteration 11130
starting iteration 11140
starting iteration 11150
starting iteration 11160
starting iteration 11170
starting iteration 11180
starting iteration 11190
starting iteration 11200
starting iteration 11210
starting iteration 11220
starting iteration 11230
starting iteration 11240
starting iteration 11250
starting iteration 11260
Finished epoch 23 / 28: cost 1.454009, train: 0.569000, val 0.512000, lr 1.32
9441e-05
starting iteration 11270
starting iteration 11280
starting iteration 11290
starting iteration 11300
starting iteration 11310
starting iteration 11320
starting iteration 11330
starting iteration 11340
starting iteration 11350
starting iteration 11360
starting iteration 11370
starting iteration 11380
starting iteration 11390
starting iteration 11400
starting iteration 11410
starting iteration 11420
starting iteration 11430
starting iteration 11440
starting iteration 11450
starting iteration 11460
starting iteration 11470
starting iteration 11480
```

starting iteration 11490  
starting iteration 11500  
starting iteration 11510  
starting iteration 11520  
starting iteration 11530  
starting iteration 11540  
starting iteration 11550  
starting iteration 11560  
starting iteration 11570  
starting iteration 11580  
starting iteration 11590  
starting iteration 11600  
starting iteration 11610  
starting iteration 11620  
starting iteration 11630  
starting iteration 11640  
starting iteration 11650  
starting iteration 11660  
starting iteration 11670  
starting iteration 11680  
starting iteration 11690  
starting iteration 11700  
starting iteration 11710  
starting iteration 11720  
starting iteration 11730  
starting iteration 11740  
starting iteration 11750  
Finished epoch 24 / 28: cost 1.509655, train: 0.538000, val 0.520000, lr 1.19  
6497e-05  
starting iteration 11760  
starting iteration 11770  
starting iteration 11780  
starting iteration 11790  
starting iteration 11800  
starting iteration 11810  
starting iteration 11820  
starting iteration 11830  
starting iteration 11840  
starting iteration 11850  
starting iteration 11860  
starting iteration 11870  
starting iteration 11880  
starting iteration 11890  
starting iteration 11900  
starting iteration 11910  
starting iteration 11920  
starting iteration 11930  
starting iteration 11940  
starting iteration 11950  
starting iteration 11960  
starting iteration 11970  
starting iteration 11980  
starting iteration 11990  
starting iteration 12000  
starting iteration 12010  
starting iteration 12020  
starting iteration 12030

```
starting iteration 12040
starting iteration 12050
starting iteration 12060
starting iteration 12070
starting iteration 12080
starting iteration 12090
starting iteration 12100
starting iteration 12110
starting iteration 12120
starting iteration 12130
starting iteration 12140
starting iteration 12150
starting iteration 12160
starting iteration 12170
starting iteration 12180
starting iteration 12190
starting iteration 12200
starting iteration 12210
starting iteration 12220
starting iteration 12230
starting iteration 12240
Finished epoch 25 / 28: cost 1.430350, train: 0.575000, val 0.518000, lr 1.07
6847e-05
starting iteration 12250
starting iteration 12260
starting iteration 12270
starting iteration 12280
starting iteration 12290
starting iteration 12300
starting iteration 12310
starting iteration 12320
starting iteration 12330
starting iteration 12340
starting iteration 12350
starting iteration 12360
starting iteration 12370
starting iteration 12380
starting iteration 12390
starting iteration 12400
starting iteration 12410
starting iteration 12420
starting iteration 12430
starting iteration 12440
starting iteration 12450
starting iteration 12460
starting iteration 12470
starting iteration 12480
starting iteration 12490
starting iteration 12500
starting iteration 12510
starting iteration 12520
starting iteration 12530
starting iteration 12540
starting iteration 12550
starting iteration 12560
starting iteration 12570
starting iteration 12580
```

```
starting iteration 12590
starting iteration 12600
starting iteration 12610
starting iteration 12620
starting iteration 12630
starting iteration 12640
starting iteration 12650
starting iteration 12660
starting iteration 12670
starting iteration 12680
starting iteration 12690
starting iteration 12700
starting iteration 12710
starting iteration 12720
starting iteration 12730
Finished epoch 26 / 28: cost 1.519867, train: 0.580000, val 0.515000, lr 9.69
1623e-06
starting iteration 12740
starting iteration 12750
starting iteration 12760
starting iteration 12770
starting iteration 12780
starting iteration 12790
starting iteration 12800
starting iteration 12810
starting iteration 12820
starting iteration 12830
starting iteration 12840
starting iteration 12850
starting iteration 12860
starting iteration 12870
starting iteration 12880
starting iteration 12890
starting iteration 12900
starting iteration 12910
starting iteration 12920
starting iteration 12930
starting iteration 12940
starting iteration 12950
starting iteration 12960
starting iteration 12970
starting iteration 12980
starting iteration 12990
starting iteration 13000
starting iteration 13010
starting iteration 13020
starting iteration 13030
starting iteration 13040
starting iteration 13050
starting iteration 13060
starting iteration 13070
starting iteration 13080
starting iteration 13090
starting iteration 13100
starting iteration 13110
starting iteration 13120
starting iteration 13130
```

```
starting iteration 13140
starting iteration 13150
starting iteration 13160
starting iteration 13170
starting iteration 13180
starting iteration 13190
starting iteration 13200
starting iteration 13210
starting iteration 13220
Finished epoch 27 / 28: cost 1.428485, train: 0.555000, val 0.521000, lr 8.72
2461e-06
starting iteration 13230
starting iteration 13240
starting iteration 13250
starting iteration 13260
starting iteration 13270
starting iteration 13280
starting iteration 13290
starting iteration 13300
starting iteration 13310
starting iteration 13320
starting iteration 13330
starting iteration 13340
starting iteration 13350
starting iteration 13360
starting iteration 13370
starting iteration 13380
starting iteration 13390
starting iteration 13400
starting iteration 13410
starting iteration 13420
starting iteration 13430
starting iteration 13440
starting iteration 13450
starting iteration 13460
starting iteration 13470
starting iteration 13480
starting iteration 13490
starting iteration 13500
starting iteration 13510
starting iteration 13520
starting iteration 13530
starting iteration 13540
starting iteration 13550
starting iteration 13560
starting iteration 13570
starting iteration 13580
starting iteration 13590
starting iteration 13600
starting iteration 13610
starting iteration 13620
starting iteration 13630
starting iteration 13640
starting iteration 13650
starting iteration 13660
starting iteration 13670
starting iteration 13680
```

```


starting iteration 13690
starting iteration 13700
starting iteration 13710
Finished epoch 28 / 28: cost 1.380179, train: 0.551000, val 0.533000, lr 7.85
0214e-06
finished optimization. best validation accuracy: 0.533000

```

```

In [16]: # visualize the weights
show_net_weights(best_model)

```



## Run on the test set

When you are done experimenting, you should evaluate your final trained network on the test set.

**We will give you extra bonus point for every 1% of accuracy above 56%.**

```

In [17]: scores_test = two_layer_net(X_test, best_model)
print 'Test accuracy: ', np.mean(np.argmax(scores_test, axis=1) == y_test)

Test accuracy: 0.52

```

```

In [ ]:

```

## Modular neural nets

In the previous exercise, we computed the loss and gradient for a two-layer neural network in a single monolithic function. This isn't very difficult for a small two-layer network, but would be tedious and error-prone for larger networks. Ideally we want to build networks using a more modular design so that we can snap together different types of layers and loss functions in order to quickly experiment with different architectures.

In this exercise we will implement this approach, and develop a number of different layer types in isolation that can then be easily plugged together. For each layer we will implement forward and backward functions. The forward function will receive data, weights, and other parameters, and will return both an output and a cache object that stores data needed for the backward pass. The backward function will receive upstream derivatives and the cache object, and will return gradients with respect to the data and all of the weights. This will allow us to write code that looks like this:

```
def two_layer_net(X, W1, b1, W2, b2, reg):
    # Forward pass; compute scores
    s1, fc1_cache = affine_forward(X, W1, b1)
    a1, relu_cache = relu_forward(s1)
    scores, fc2_cache = affine_forward(a1, W2, b2)

    # Loss functions return data loss and gradients on scores
    data_loss, dscores = svm_loss(scores, y)

    # Compute backward pass
    da1, dW2, db2 = affine_backward(dscores, fc2_cache)
    ds1 = relu_backward(da1, relu_cache)
    dX, dW1, db1 = affine_backward(ds1, fc1_cache)

    # A real network would add regularization here

    # Return loss and gradients
    return loss, dW1, db1, dW2, db2
```



```
In [3]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from cs231n.gradient_check import eval_numerical_gradient_array, eval_numerical_gradient
from cs231n.layers import *

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

## Affine layer: forward

Open the file `cs231n/layers.py` and implement the `affine_forward` function.

Once you are done we will test your can test your implementation by running the following:

```
In [5]: # Test the affine_forward function

num_inputs = 2
input_shape = (4, 5, 6)
output_dim = 3

input_size = num_inputs * np.prod(input_shape)
weight_size = output_dim * np.prod(input_shape)

x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape), output_dim)
b = np.linspace(-0.3, 0.1, num=output_dim)

out, _ = affine_forward(x, w, b)
correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
                        [ 3.25553199,  3.5141327,  3.77273342]])

# Compare your output with ours. The error should be around 1e-9.
print 'Testing affine_forward function:'
print 'difference: ', rel_error(out, correct_out)

Testing affine_forward function:
difference: 9.76984772881e-10
```

## Affine layer: backward

Now implement the `affine_backward` function. You can test your implementation using numeric gradient checking.

```
In [6]: # Test the affine_backward function

x = np.random.randn(10, 2, 3)
w = np.random.randn(6, 5)
b = np.random.randn(5)
dout = np.random.randn(10, 5)

dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0],
x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0],
w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0],
b, dout)

_, cache = affine_forward(x, w, b)
dx, dw, db = affine_backward(dout, cache)

# The error should be less than 1e-10
print 'Testing affine_backward function:'
print 'dx error: ', rel_error(dx_num, dx)
print 'dw error: ', rel_error(dw_num, dw)
print 'db error: ', rel_error(db_num, db)

Testing affine_backward function:
dx error:  6.99811149861e-10
dw error:  9.14830163174e-11
db error:  2.16251754573e-11
```

## ReLU layer: forward

Implement the relu\_forward function and test your implementation by running the following:

```
In [7]: # Test the relu_forward function

x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)

out, _ = relu_forward(x)
correct_out = np.array([[ 0.,          0.,          0.,          0.],
                        [ 0.,          0.,          0.04545455, 0.13636364],
                        [ 0.22727273, 0.31818182, 0.40909091, 0.5,
]])

# Compare your output with ours. The error should be around 1e-8
print 'Testing relu_forward function:'
print 'difference: ', rel_error(out, correct_out)

Testing relu_forward function:
difference:  4.99999979802e-08
```

## ReLU layer: backward

Implement the `relu_backward` function and test your implementation using numeric gradient checking:

```
In [8]: x = np.random.randn(10, 10)
dout = np.random.randn(*x.shape)

dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)

_, cache = relu_forward(x)
dx = relu_backward(dout, cache)

# The error should be around 1e-12
print 'Testing relu_backward function:'
print 'dx error: ', rel_error(dx_num, dx)

Testing relu_backward function:
dx error:  3.27564222471e-12
```

## Loss layers: Softmax and SVM

You implemented these loss functions in the last assignment, so we'll give them to you for free here. It's still a good idea to test them to make sure they work correctly.

```
In [9]: num_classes, num_inputs = 10, 50
x = 0.001 * np.random.randn(num_inputs, num_classes)
y = np.random.randint(num_classes, size=num_inputs)

dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x,
verbose=False)
loss, dx = svm_loss(x, y)

# Test svm_loss function. Loss should be around 9 and dx error should be 1e-9
print 'Testing svm_loss:'
print 'loss: ', loss
print 'dx error: ', rel_error(dx_num, dx)

dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x, verbose=False)
loss, dx = softmax_loss(x, y)

# Test softmax_loss function. Loss should be 2.3 and dx error should be 1e-8
print '\nTesting softmax_loss:'
print 'loss: ', loss
print 'dx error: ', rel_error(dx_num, dx)

Testing svm_loss:
loss: 8.99999185188
dx error: 8.18289447289e-10

Testing softmax_loss:
loss: 2.30258473372
dx error: 8.68969367873e-09
```

Listing 7.1: *neural\_net.py*

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3
4  def init_two_layer_model(input_size, hidden_size, output_size):
5      """
6      Initialize the weights and biases for a two-layer fully connected neural
7      network. The net has an input dimension of D, a hidden layer dimension of H,
8      and performs classification over C classes. Weights are initialized to small
9      random values and biases are initialized to zero.
10
11      Inputs:
12      - input_size: The dimension D of the input data
13      - hidden_size: The number of neurons H in the hidden layer
14      - output_size: The number of classes C
15
16      Returns:
17      A dictionary mapping parameter names to arrays of parameter values. It has
18      the following keys:
19      - W1: First layer weights; has shape (D, H)
20      - b1: First layer biases; has shape (H,)
21      - W2: Second layer weights; has shape (H, C)
22      - b2: Second layer biases; has shape (C,)
23      """
24      # initialize a model
25      model = {}
26      model['W1'] = 0.00001 * np.random.randn(input_size, hidden_size)
27      model['b1'] = np.zeros(hidden_size)
28      model['W2'] = 0.00001 * np.random.randn(hidden_size, output_size)
29      model['b2'] = np.zeros(output_size)
30      return model
31
32  def two_layer_net(X, model, y=None, reg=0.0):
33      """
34      Compute the loss and gradients for a two layer fully connected neural network.
35      The net has an input dimension of D, a hidden layer dimension of H, and
36      performs classification over C classes. We use a softmax loss function and L2
37      regularization the the weight matrices. The two layer net should use a ReLU
38      nonlinearity after the first affine layer.
39
40      The two layer net has the following architecture:
41
42      input - fully connected layer - ReLU - fully connected layer - softmax
43
44      The outputs of the second fully-connected layer are the scores for each
45      class.
46
47      Inputs:
48      - X: Input data of shape (N, D). Each X[i] is a training sample.
49      - model: Dictionary mapping parameter names to arrays of parameter values.
50      It should contain the following:
51      - W1: First layer weights; has shape (D, H)
52      - b1: First layer biases; has shape (H,)
53      - W2: Second layer weights; has shape (H, C)
54      - b2: Second layer biases; has shape (C,)
55      - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
56      an integer in the range 0 <= y[i] < C. This parameter is optional; if it
57      is not passed then we only return scores, and if it is passed then we
58      instead return the loss and gradients.
59      - reg: Regularization strength.
60
61      Returns:
62      If y not is passed, return a matrix scores of shape (N, C) where scores[i, c]
63      is the score for class c on input X[i].
64
65      If y is not passed, instead return a tuple of:
66      - loss: Loss (data loss and regularization loss) for this batch of training
67      samples.
68      - grads: Dictionary mapping parameter names to gradients of those parameters
69      with respect to the loss function. This should have the same keys as model.
70      """
71

```

```

72 # unpack variables from the model dictionary
73 W1,b1,W2,b2 = model['W1'], model['b1'], model['W2'], model['b2']
74 N, D = X.shape
75 # input - fully connected layer - ReLU - fully connected layer - softmax
76 #Activation function
77 reluF = lambda x: np.maximum(0, x)
78 #Compared to lecture notes we switch X and W1 because to multiple the inputs
79 #with the correct weights
80 h1 = reluF(np.dot(X, W1) + b1)
81 scores = np.dot(h1, W2) + b2
82 #####
83 #                               END OF YOUR CODE                               #
84 #####
85
86 # If the targets are not given then jump out, we're done
87 if y is None:
88     return scores
89
90 #####
91 # TODO: Finish the forward pass, and compute the loss. This should include #
92 # both the data loss and L2 regularization for W1 and W2. Store the result #
93 # in the variable loss, which should be a scalar. Use the Softmax #
94 # classifier loss. So that your results match ours, multiply the #
95 # regularization loss by 0.5 #
96 #####
97 # compute the loss
98 expScores = np.exp(scores)
99 rowSum = expScores.sum(axis=1, keepdims=True)
100 # Normalized scores
101 propScores = expScores / rowSum
102 logprob_correctLabel = -np.log(propScores[range(N),y])
103 softmax_loss = 1/float(N) * np.sum(logprob_correctLabel)
104 #regularization loss
105 reg_loss = 0.5 * reg * np.sum(W1*W1) + 0.5 * reg * np.sum(W2 * W2)
106 #Final loss
107 loss = softmax_loss + reg_loss
108 #####
109 #                               END OF YOUR CODE                               #
110 #####
111
112 # compute the gradients
113 grads = {}
114
115 #####
116 # TODO: Compute the backward pass, computing the derivatives of the weights #
117 # and biases. Store the results in the grads dictionary. For example, #
118 # grads['W1'] should store the gradient on W1, and be a matrix of same size #
119 #####
120 # Firstly we calculate the gradient on the scores.
121 # The gradient from the loss function is simply -1.
122 # This is subtracted from the correct scores for each
123 # dscores are the probabilities for all classes as a row for each sample
124 dscores = propScores
125 #For each row(sample) in dscores 1 is subtracted from the correct element
126 #specified by y
127 dscores[range(N),y] -= 1
128 # We then divide all elements with N(number of samples)
129 dscores /= N
130
131 # The gradient for W2 is simply the output from the RELU activation function (h1)
132 # multiplied with the dscores that contains the gradient on the scores.
133 # d/dw(w*x) = x which is our h1 then we get the input times dscores
134 grads['W2'] = np.dot(h1.T, dscores)
135 #bias is just the sum of the dscores
136 grads['b2'] = np.sum(dscores, axis=0)
137
138 # next backprop into hidden layer. This is the scores multiplied with the weights
139 # for second layer
140 dhidden = np.dot(dscores, W2.T)
141 # backprop the ReLU non-linearity.
142 #For elements < or equals 0 we set them equals to 0
143 # remember how Relu is just max, so it routes the gradients
144 dhidden[h1 <= 0] = 0

```

```

145 # same thing as second layer - d/dw(w*x) = x, so x times our gradient for dhidden
146 grads['W1'] = np.dot(X.T, dhidden)
147 grads['b1'] = np.sum(dhidden, axis=0)
148
149 # adding gradient for regularization
150 # d/dw(1/2*reg*W1*w1) = reg * W1
151 grads['W1'] += reg * W1
152 grads['W2'] += reg * W2
153 #####
154 #                                     END OF YOUR CODE                                     #
155 #####
156
157 return loss, grads

```

Listing 7.2: *classifier\_trainer.py*

```

1 import numpy as np
2
3
4 class ClassifierTrainer(object):
5     """ The trainer class performs SGD with momentum on a cost function """
6     def __init__(self):
7         self.step_cache = {} # for storing velocities in momentum update
8
9     def train(self, X, y, X_val, y_val,
10 model, loss_function,
11 reg=0.0,
12 learning_rate=1e-2, momentum=0, learning_rate_decay=0.95,
13 update='momentum', sample_batches=True,
14 num_epochs=30, batch_size=100, acc_frequency=None,
15 verbose=False):
16     """
17     Optimize the parameters of a model to minimize a loss function. We use
18     training data X and y to compute the loss and gradients, and periodically
19     check the accuracy on the validation set.
20
21     Inputs:
22     - X: Array of training data; each X[i] is a training sample.
23     - y: Vector of training labels; y[i] gives the label for X[i].
24     - X_val: Array of validation data
25     - y_val: Vector of validation labels
26     - model: Dictionary that maps parameter names to parameter values. Each
27     parameter value is a numpy array.
28     - loss_function: A function that can be called in the following ways:
29     scores = loss_function(X, model, reg=reg)
30     loss, grads = loss_function(X, model, y, reg=reg)
31     - reg: Regularization strength. This will be passed to the loss function.
32     - learning_rate: Initial learning rate to use.
33     - momentum: Parameter to use for momentum updates.
34     - learning_rate_decay: The learning rate is multiplied by this after each
35     epoch.
36     - update: The update rule to use. One of 'sgd', 'momentum', or 'rmsprop'.
37     - sample_batches: If True, use a minibatch of data for each parameter update
38     (stochastic gradient descent); if False, use the entire training set for
39     each parameter update (gradient descent).
40     - num_epochs: The number of epochs to take over the training data.
41     - batch_size: The number of training samples to use at each iteration.
42     - acc_frequency: If set to an integer, we compute the training and
43     validation set error after every acc_frequency iterations.
44     - verbose: If True, print status after each epoch.
45
46     Returns a tuple of:
47     - best_model: The model that got the highest validation accuracy during
48     training.
49     - loss_history: List containing the value of the loss function at each
50     iteration.
51     - train_acc_history: List storing the training set accuracy at each epoch.
52     - val_acc_history: List storing the validation set accuracy at each epoch.
53     """
54

```



```

55 N = X.shape[0]
56
57 if sample_batches:
58     iterations_per_epoch = N / batch_size # using SGD
59 else:
60     iterations_per_epoch = 1 # using GD
61 num_iters = num_epochs * iterations_per_epoch
62 epoch = 0
63 best_val_acc = 0.0
64 best_model = {}
65 loss_history = []
66 train_acc_history = []
67 val_acc_history = []
68 for it in xrange(num_iters):
69     if it % 10 == 0: print 'starting iteration ', it
70
71     # get batch of data
72     if sample_batches:
73         batch_mask = np.random.choice(N, batch_size)
74         X_batch = X[batch_mask]
75         y_batch = y[batch_mask]
76     else:
77         # no SGD used, full gradient descent
78         X_batch = X
79         y_batch = y
80
81     # evaluate cost and gradient
82     cost, grads = loss_function(X_batch, model, y_batch, reg)
83     loss_history.append(cost)
84
85     # perform a parameter update
86     for p in model:
87         # compute the parameter step
88         if update == 'sgd':
89             dx = -learning_rate * grads[p]
90         elif update == 'momentum':
91             if not p in self.step_cache:
92                 self.step_cache[p] = np.zeros(grads[p].shape)
93             dx = np.zeros_like(grads[p]) # you can remove this after
94             #####
95             # TODO: implement the momentum update formula and store the step #
96             # update into variable dx. You should use the variable #
97             # step_cache[p] and the momentum strength is stored in momentum. #
98             # Don't forget to also update the step_cache[p]. #
99             #####
100             # Momentum update
101             # Momentum strength is a coefficient explains to friction,
102             # moment strength damps velocity and reduces the kinetic energy
103             # it ensures that the particle stops at the bottom
104             # step_cache is our velocity at time t-1
105             # From that we subtract the learning rate multiplied the gradients,
106             # because we go in the opposite direction of the gradient
107             dx = momentum * self.step_cache[p] - learning_rate * grads[p] # integrate velocity
108             self.step_cache[p] = dx
109
110             #####
111             #                               END OF YOUR CODE                               #
112             #####
113         elif update == 'rmsprop':
114             decay_rate = 0.99 # you could also make this an option
115             if not p in self.step_cache:
116                 self.step_cache[p] = np.zeros(grads[p].shape)
117             #####
118             # TODO: implement the RMSProp update and store the parameter update #
119             # dx. Don't forget to also update step_cache[p]. Use smoothing 1e-8 #
120             #####
121             #eps = 1e-8
122             eps = 1e5
123             self.step_cache[p] = decay_rate * self.step_cache[p] + (1 - decay_rate) * grads[p]**2
124             dx = - learning_rate * grads[p] / (np.sqrt(self.step_cache[p]) + eps)
125
126
127     #####

```

```

128 #                                     END OF YOUR CODE                                     #
129 #####
130 else:
131     raise ValueError('Unrecognized update type "%s" ' % update)
132
133 # update the parameters
134 model[p] += dx
135
136 # every epoch perform an evaluation on the validation set
137 first_it = (it == 0)
138 epoch_end = (it + 1) % iterations_per_epoch == 0
139 acc_check = (acc.frequency is not None and it % acc.frequency == 0)
140 if first_it or epoch_end or acc_check:
141     if it > 0 and epoch_end:
142         # decay the learning rate
143         learning_rate *= learning_rate_decay
144         epoch += 1
145
146 # evaluate train accuracy
147 if N > 1000:
148     train_mask = np.random.choice(N, 1000)
149     X_train_subset = X[train_mask]
150     y_train_subset = y[train_mask]
151 else:
152     X_train_subset = X
153     y_train_subset = y
154 scores_train = loss_function(X_train_subset, model)
155 y_pred_train = np.argmax(scores_train, axis=1)
156 train_acc = np.mean(y_pred_train == y_train_subset)
157 train_acc_history.append(train_acc)
158
159 # evaluate val accuracy
160 scores_val = loss_function(X_val, model)
161 y_pred_val = np.argmax(scores_val, axis=1)
162 val_acc = np.mean(y_pred_val == y_val)
163 val_acc_history.append(val_acc)
164
165 # keep track of the best model based on validation accuracy
166 if val_acc > best_val_acc:
167     # make a copy of the model
168     best_val_acc = val_acc
169     best_model = {}
170     for p in model:
171         best_model[p] = model[p].copy()
172
173 # print progress if needed
174 if verbose:
175     print ('Finished epoch %d / %d: cost %f, train: %f, val %f, lr %e'
176           % (epoch, num_epochs, cost, train_acc, val_acc, learning_rate))
177
178 if verbose:
179     print 'finished optimization. best validation accuracy: %f' % (best_val_acc, )
180 # return the best model and the training history statistics
181 return best_model, loss_history, train_acc_history, val_acc_history

```

Listing 7.3: *layers.py*

```

1 import numpy as np
2
3 def affine_forward(x, w, b):
4     """
5     Computes the forward pass for an affine (fully-connected) layer.
6
7     The input x has shape (N, d_1, ..., d_k) where x[i] is the ith input.
8     We multiply this against a weight matrix of shape (D, M) where
9     D = \prod_i d_i
10
11     Inputs:
12     x - Input data, of shape (N, d_1, ..., d_k)
13     w - Weights, of shape (D, M)

```

```

14 b - Biases, of shape (M,)
15
16 Returns a tuple of:
17 - out: output, of shape (N, M)
18 - cache: (x, w, b)
19 """
20 out = None
21 #####
22 # TODO: Implement the affine forward pass. Store the result in out. You #
23 # will need to reshape the input into rows. #
24 #####
25 # First we reshape X to multiply it with the incoming weights
26 # We get column and row size and then reshape
27 row_size = x.shape[0]
28 col_size = np.prod(x.shape[1:])
29 x.reshape = x.reshape(row_size, col_size)
30 # To execute the forward pass we simply need to multiply the inputs with the weights
31 out = np.dot(x.reshape, w) + b
32 #####
33 #                               END OF YOUR CODE                               #
34 #####
35 cache = (x, w, b)
36 return out, cache
37
38
39 def affine_backward(dout, cache):
40     """
41     Computes the backward pass for an affine layer.
42
43     Inputs:
44     - dout: Upstream derivative, of shape (N, M)
45     - cache: Tuple of:
46     - x: Input data, of shape (N, d_1, ..., d_k)
47     - w: Weights, of shape (D, M)
48
49     Returns a tuple of:
50     - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
51     - dw: Gradient with respect to w, of shape (D, M)
52     - db: Gradient with respect to b, of shape (M,)
53     """
54     x, w, b = cache
55     dx, dw, db = None, None, None
56     #####
57     # TODO: Implement the affine backward pass. #
58     #####
59     # First we reshape X to multiply it with the incoming weights
60     # We get column and row size and then reshape
61     row_size = x.shape[0]
62     col_size = np.prod(x.shape[1:])
63     x.reshape = x.reshape(row_size, col_size)
64
65     # After reshaping we calculate the backward pass
66
67     # Gradient with respect to x
68     # Gradient of x*w with respect to x is simply w,
69     # We multiply that with the upstream gradient
70     dx2 = np.dot(dout, w.T)
71     dx = np.reshape(dx2, x.shape)
72
73     # Gradient with respect to weights
74     # Gradient of x*w with respect to w is simply x,
75     # We multiply that with the upstream gradient
76     dw = np.dot(x.reshape.T, dout)
77
78     # Gradient with respect to bias
79     # Biases are added so the gradient is simply 1,
80     # We multiply that with the upstream gradient.
81     db = np.dot(dout.T, np.ones(row_size))
82
83     #####
84     #                               END OF YOUR CODE                               #
85     #####
86     return dx, dw, db

```

```

87
88
89 def relu_forward(x):
90     """
91     Computes the forward pass for a layer of rectified linear units (ReLU).
92
93     Input:
94     - x: Inputs, of any shape
95
96     Returns a tuple of:
97     - out: Output, of the same shape as x
98     - cache: x
99     """
100     out = None
101     #####
102     # TODO: Implement the ReLU forward pass. #
103     #####
104     reluF = lambda x: np.maximum(0, x)
105     out = reluF(x)
106     #####
107     #                               END OF YOUR CODE                               #
108     #####
109     cache = x
110     return out, cache
111
112
113 def relu_backward(dout, cache):
114     """
115     Computes the backward pass for a layer of rectified linear units (ReLU).
116
117     Input:
118     - dout: Upstream derivatives, of any shape
119     - cache: Input x, of same shape as dout
120
121     Returns:
122     - dx: Gradient with respect to x
123     """
124     dx, x = None, cache
125     #####
126     # TODO: Implement the ReLU backward pass. #
127     #####
128     #reluf function
129     reluF = lambda x: np.maximum(0, x)
130
131     out = reluF(x)
132     # Reluf is a max gate and so we can think of it as a router of gradients
133     # the max value is the one that the gradient is routated to
134     # we simply set the out value to 1 if the out value is bigger than 0
135     out[out > 0] = 1
136
137     # Multiply out with upstream gradient, to "route" the gradient
138     dx = out * dout
139     #####
140     #                               END OF YOUR CODE                               #
141     #####
142     return dx
143
144 def dropout_forward(x, dropout_param):
145     """
146     Performs the forward pass for (inverted) dropout.
147
148     Inputs:
149     - x: Input data, of any shape
150     - dropout_param: A dictionary with the following keys:
151     - p: Dropout parameter. We keep each neuron output with probability p.
152     - mode: 'test' or 'train'. If the mode is train, then perform dropout;
153     if the mode is test, then just return the input.
154     - seed: Seed for the random number generator. Passing seed makes this
155     function deterministic, which is needed for gradient checking but not in
156     real networks.
157
158     Outputs:
159     - out: Array of the same shape as x.

```

```

160 - cache: A tuple (dropout_param, mask). In training mode, mask is the dropout
161 mask that was used to multiply the input; in test mode, mask is None.
162 """
163 p, mode = dropout_param['p'], dropout_param['mode']
164 if 'seed' in dropout_param:
165     np.random.seed(dropout_param['seed'])
166
167     mask = None
168     out = None
169
170     if mode == 'train':
171         #####
172         # TODO: Implement the training phase forward pass for inverted dropout. #
173         # Store the dropout mask in the mask variable. #
174         #####
175         pass
176         #####
177         #                               END OF YOUR CODE                               #
178         #####
179     elif mode == 'test':
180         #####
181         # TODO: Implement the test phase forward pass for inverted dropout. #
182         #####
183         pass
184         #####
185         #                               END OF YOUR CODE                               #
186         #####
187
188     cache = (dropout_param, mask)
189     out = out.astype(x.dtype, copy=False)
190
191     return out, cache
192
193
194 def dropout_backward(dout, cache):
195     """
196     Perform the backward pass for (inverted) dropout.
197
198     Inputs:
199     - dout: Upstream derivatives, of any shape
200     - cache: (dropout_param, mask) from dropout_forward.
201     """
202     dropout_param, mask = cache
203     mode = dropout_param['mode']
204     if mode == 'train':
205         #####
206         # TODO: Implement the training phase forward pass for inverted dropout. #
207         # Store the dropout mask in the mask variable. #
208         #####
209         pass
210         #####
211         #                               END OF YOUR CODE                               #
212         #####
213     elif mode == 'test':
214         dx = dout
215         return dx
216
217
218 def conv_forward_naive(x, w, b, conv_param):
219     """
220     A naive implementation of the forward pass for a convolutional layer.
221
222     The input consists of N data points, each with C channels, height H and width
223     W. We convolve each input with F different filters, where each filter spans
224     all C channels and has height HH and width HH.
225
226     Input:
227     - x: Input data of shape (N, C, H, W)
228     - w: Filter weights of shape (F, C, HH, WW)
229     - b: Biases, of shape (F,)
230     - conv_param: A dictionary with the following keys:
231     - 'stride': The number of pixels between adjacent receptive fields in the
232       horizontal and vertical directions.

```

```

233 - 'pad': The number of pixels that will be used to zero-pad the input.
234
235 Returns a tuple of:
236 - out: Output data, of shape (N, F, H', W') where H' and W' are given by
237    $H' = 1 + (H + 2 * \text{pad} - \text{HH}) / \text{stride}$ 
238    $W' = 1 + (W + 2 * \text{pad} - \text{WW}) / \text{stride}$ 
239 - cache: (x, w, b, conv_param)
240 """
241 out = None
242 #####
243 # TODO: Implement the convolutional forward pass. #
244 # Hint: you can use the function np.pad for padding. #
245 #####
246 pass
247 #####
248 #                               END OF YOUR CODE                               #
249 #####
250 cache = (x, w, b, conv_param)
251 return out, cache
252
253
254 def conv_backward_naive(dout, cache):
255     """
256     A naive implementation of the backward pass for a convolutional layer.
257
258     Inputs:
259     - dout: Upstream derivatives.
260     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
261
262     Returns a tuple of:
263     - dx: Gradient with respect to x
264     - dw: Gradient with respect to w
265     - db: Gradient with respect to b
266     """
267     dx, dw, db = None, None, None
268     #####
269     # TODO: Implement the convolutional backward pass. #
270     #####
271     pass
272     #####
273     #                               END OF YOUR CODE                               #
274     #####
275     return dx, dw, db
276
277
278 def max_pool_forward_naive(x, pool_param):
279     """
280     A naive implementation of the forward pass for a max pooling layer.
281
282     Inputs:
283     - x: Input data, of shape (N, C, H, W)
284     - pool_param: dictionary with the following keys:
285       - 'pool_height': The height of each pooling region
286       - 'pool_width': The width of each pooling region
287       - 'stride': The distance between adjacent pooling regions
288
289     Returns a tuple of:
290     - out: Output data
291     - cache: (x, pool_param)
292     """
293     out = None
294     #####
295     # TODO: Implement the max pooling forward pass #
296     #####
297     pass
298     #####
299     #                               END OF YOUR CODE                               #
300     #####
301     cache = (x, pool_param)
302     return out, cache
303
304
305 def max_pool_backward_naive(dout, cache):

```

```

306 """
307 A naive implementation of the backward pass for a max pooling layer.
308
309 Inputs:
310 - dout: Upstream derivatives
311 - cache: A tuple of (x, pool_param) as in the forward pass.
312
313 Returns:
314 - dx: Gradient with respect to x
315 """
316 dx = None
317 #####
318 # TODO: Implement the max pooling backward pass #
319 #####
320 pass
321 #####
322 #                                     END OF YOUR CODE #
323 #####
324 return dx
325
326
327 def svm_loss(x, y):
328 """
329 Computes the loss and gradient using for multiclass SVM classification.
330
331 Inputs:
332 - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
333   for the ith input.
334 - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
335   0 <= y[i] < C
336
337 Returns a tuple of:
338 - loss: Scalar giving the loss
339 - dx: Gradient of the loss with respect to x
340 """
341 N = x.shape[0]
342 correct_class_scores = x[np.arange(N), y]
343 margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
344 margins[np.arange(N), y] = 0
345 loss = np.sum(margins) / N
346 num_pos = np.sum(margins > 0, axis=1)
347 dx = np.zeros_like(x)
348 dx[margins > 0] = 1
349 dx[np.arange(N), y] -= num_pos
350 dx /= N
351 return loss, dx
352
353
354 def softmax_loss(x, y):
355 """
356 Computes the loss and gradient for softmax classification.
357
358 Inputs:
359 - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
360   for the ith input.
361 - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
362   0 <= y[i] < C
363
364 Returns a tuple of:
365 - loss: Scalar giving the loss
366 - dx: Gradient of the loss with respect to x
367 """
368 probs = np.exp(x - np.max(x, axis=1, keepdims=True))
369 probs /= np.sum(probs, axis=1, keepdims=True)
370 N = x.shape[0]
371 loss = -np.sum(np.log(probs[np.arange(N), y])) / N
372 dx = probs.copy()
373 dx[np.arange(N), y] -= 1
374 dx /= N
375 return loss, dx

```

## 8 Exercise 2



## Modular neural nets

In the previous exercise, we started to build modules/general layers for implementing large neural networks. In this exercise, we will expand on this by implementing a convolutional layer, max pooling layer and a dropout layer. For each layer we will implement forward and backward functions. The forward function will receive data, weights, and other parameters, and will return both an output and a cache object that stores data needed for the backward pass. The backward function will receive upstream derivatives and the cache object, and will return gradients with respect to the data and all of the weights. This will allow us to write code that looks like this:

```
def two_layer_net(X, W1, b1, W2, b2, reg):
    # Forward pass; compute scores
    s1, fc1_cache = affine_forward(X, W1, b1)
    a1, relu_cache = relu_forward(s1)
    scores, fc2_cache = affine_forward(a1, W2, b2)

    # Loss functions return data loss and gradients on scores
    data_loss, dscores = svm_loss(scores, y)

    # Compute backward pass
    da1, dW2, db2 = affine_backward(dscores, fc2_cache)
    ds1 = relu_backward(da1, relu_cache)
    dX, dW1, db1 = affine_backward(ds1, fc1_cache)

    # A real network would add regularization here

    # Return loss and gradients
    return loss, dW1, db1, dW2, db2
```

```
In [1]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from cs231n.gradient_check import eval_numerical_gradient_array, eval_numerical_gradient
from cs231n.layers import *

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

## Dropout layer: forward

Open the file `cs231n/layers.py` and implement the `dropout_forward` function. You should implement **inverted dropout** rather than regular dropout. We can check the forward pass by looking at the statistics of the outputs in train and test modes.

```
In [2]: # Check the dropout forward pass

x = np.random.randn(100, 100)
dropout_param_train = {'p': 0.25, 'mode': 'train'}
dropout_param_test = {'p': 0.25, 'mode': 'test'}

out_train, _ = dropout_forward(x, dropout_param_train)
out_test, _ = dropout_forward(x, dropout_param_test)

# Test dropout training mode; about 25% of the elements should be nonzero
print np.mean(out_train != 0) # expected to be ~0.25

# Test dropout test mode; all of the elements should be nonzero
print np.mean(out_test != 0) # expected to be = 1

0.2511
1.0
```

## Dropout layer: backward

Open the file `cs231n/layers.py` and implement the `dropout_backward` function. We can check the backward pass using numerical gradient checking.

```
In [3]: from cs231n.gradient_check import eval_numerical_gradient_array

        # Check the dropout backward pass

        x = np.random.randn(5, 4)
        dout = np.random.randn(*x.shape)
        dropout_param = {'p': 0.8, 'mode': 'train', 'seed': 123}

        dx_num = eval_numerical_gradient_array(lambda x: dropout_forward(x, dropout_param)[0], x, dout)

        _, cache = dropout_forward(x, dropout_param)
        dx = dropout_backward(dout, cache)

        # The error should be around 1e-12
        print 'Testing dropout_backward function:'
        print 'dx error: ', rel_error(dx_num, dx)

Testing dropout_backward function:
dx error:  3.38580207168e-12
```

## Convolution layer: forward naive

We are now ready to implement the forward pass for a convolutional layer. Implement the function `conv_forward_naive` in the file `cs231n/layers.py`.

You don't have to worry too much about efficiency at this point; just write the code in whatever way you find most clear.

You can test your implementation by running the following:

```
In [4]: x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [[ 0.21027089,  0.21661097],
                           [ 0.22847626,  0.23004637]],
                          [[ 0.50813986,  0.54309974],
                           [ 0.64082444,  0.67101435]]],
                         [[[-0.98053589, -1.03143541],
                           [-1.19128892, -1.24695841]],
                          [[ 0.69108355,  0.66880383],
                           [ 0.59480972,  0.56776003]],
                          [[ 2.36270298,  2.36904306],
                           [ 2.38090835,  2.38247847]]]])

# Compare your output to ours; difference should be around 1e-8
print 'Testing conv_forward_naive'
print 'difference: ', rel_error(out, correct_out)
```

```
Testing conv_forward_naive
difference:  2.21214764175e-08
```

## Aside: Image processing via convolutions

As fun way to both check your implementation and gain a better understanding of the type of operation that convolutional layers can perform, we will set up an input containing two images and manually set up filters that perform common image processing operations (grayscale conversion and edge detection). The convolution forward pass will apply these operations to each of the input images. We can then visualize the results as a sanity check.

```
In [5]: from scipy.misc import imread, imresize

kitten, puppy = imread('kitten.jpg'), imread('puppy.jpg')
# kitten is wide, and puppy is already square
d = kitten.shape[1] - kitten.shape[0]
kitten_cropped = kitten[:, d/2:-d/2, :]

img_size = 200 # Make this smaller if it runs too slow
x = np.zeros((2, 3, img_size, img_size))
x[0, :, :, :] = imresize(puppy, (img_size, img_size)).transpose((2, 0, 1))
x[1, :, :, :] = imresize(kitten_cropped, (img_size, img_size)).transpose((2,
0, 1))

# Set up a convolutional weights holding 2 filters, each 3x3
w = np.zeros((2, 3, 3, 3))

# The first filter converts the image to grayscale.
```

```

# Set up the red, green, and blue channels of the filter.
w[0, 0, :, :] = [[0, 0, 0], [0, 0.3, 0], [0, 0, 0]]
w[0, 1, :, :] = [[0, 0, 0], [0, 0.6, 0], [0, 0, 0]]
w[0, 2, :, :] = [[0, 0, 0], [0, 0.1, 0], [0, 0, 0]]

# Second filter detects horizontal edges in the blue channel.
w[1, 2, :, :] = [[1, 2, 1], [0, 0, 0], [-1, -2, -1]]

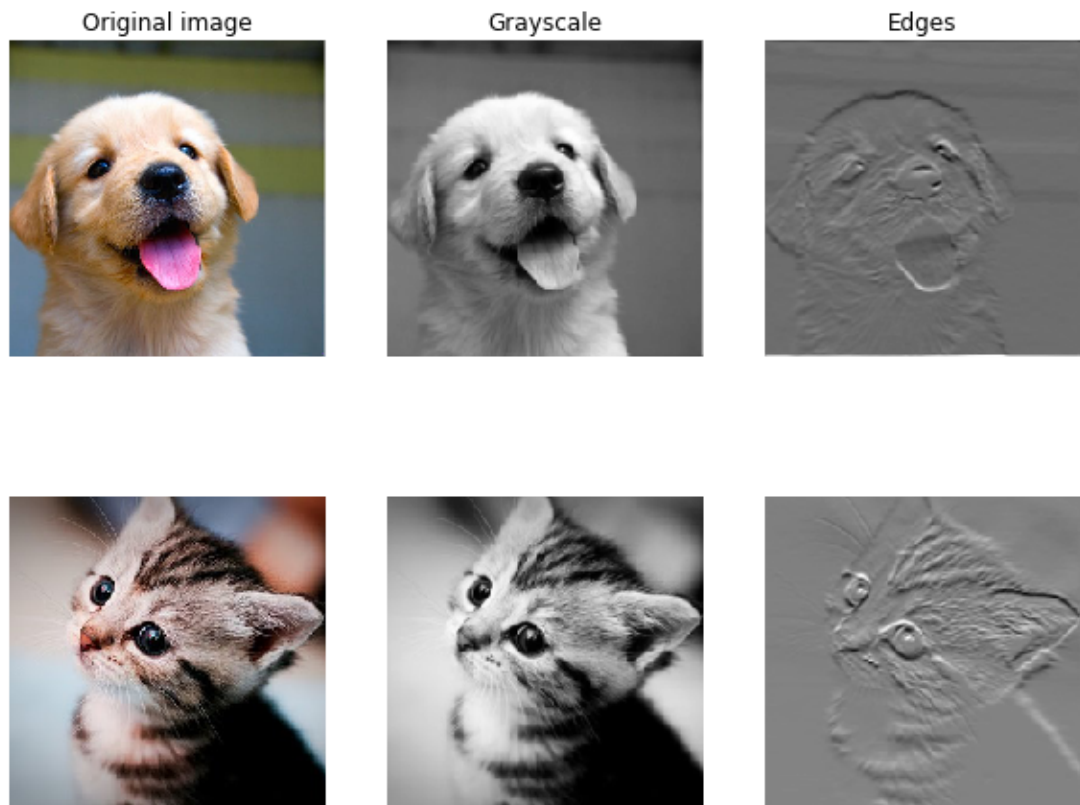
# Vector of biases. We don't need any bias for the grayscale
# filter, but for the edge detection filter we want to add 128
# to each output so that nothing is negative.
b = np.array([0, 128])

# Compute the result of convolving each input in x with each filter in w,
# offsetting by b, and storing the results in out.
out, _ = conv_forward_naive(x, w, b, {'stride': 1, 'pad': 1})

def imshow_noax(img, normalize=True):
    """ Tiny helper to show images as uint8 and remove axis labels """
    if normalize:
        img_max, img_min = np.max(img), np.min(img)
        img = 255.0 * (img - img_min) / (img_max - img_min)
    plt.imshow(img.astype('uint8'))
    plt.gca().axis('off')

# Show the original images and the results of the conv operation
plt.subplot(2, 3, 1)
imshow_noax(puppy, normalize=False)
plt.title('Original image')
plt.subplot(2, 3, 2)
imshow_noax(out[0, 0])
plt.title('Grayscale')
plt.subplot(2, 3, 3)
imshow_noax(out[0, 1])
plt.title('Edges')
plt.subplot(2, 3, 4)
imshow_noax(kitten_cropped, normalize=False)
plt.subplot(2, 3, 5)
imshow_noax(out[1, 0])
plt.subplot(2, 3, 6)
imshow_noax(out[1, 1])
plt.show()

```



## Convolution layer: backward naive

Next you need to implement the function `conv_backward_naive` in the file `cs231n/layers.py`. As usual, we will check your implementation with numeric gradient checking.

```
In [6]: x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
conv_param = {'stride': 1, 'pad': 1}

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b, c
onv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b, c
onv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b, c
onv_param)[0], b, dout)

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9
print 'Testing conv_backward_naive function'
print 'dx error: ', rel_error(dx, dx_num)
print 'dw error: ', rel_error(dw, dw_num)
print 'db error: ', rel_error(db, db_num)

Testing conv_backward_naive function
dx error:  1.42918515583e-09
dw error:  3.24894092112e-10
db error:  4.93456965076e-12
```

## Max pooling layer: forward naive

The last layer we need for a basic convolutional neural network is the max pooling layer. First implement the forward pass in the function `max_pool_forward_naive` in the file `cs231n/layers.py`.



```
In [7]: x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)
correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                          [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]],
                          [[-0.02736842, -0.01263158],
                           [ 0.03157895,  0.04631579]]],
                        [[[ 0.09052632,  0.10526316],
                           [ 0.14947368,  0.16421053]],
                          [[ 0.20842105,  0.22315789],
                           [ 0.26736842,  0.28210526]],
                          [[ 0.32631579,  0.34105263],
                           [ 0.38526316,  0.4          ]]]]])

# Compare your output with ours. Difference should be around 1e-8.
print 'Testing max_pool_forward_naive function:'
print 'difference: ', rel_error(out, correct_out)
```

```
Testing max_pool_forward_naive function:
difference: 4.16666651573e-08
```

## Max pooling layer: backward naive

Implement the backward pass for a max pooling layer in the function `max_pool_backward_naive` in the file `cs231n/layers.py`. As always we check the correctness of the backward pass using numerical gradient checking.

```
In [8]: x = np.random.randn(3, 2, 8, 8)
dout = np.random.randn(3, 2, 4, 4)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x, pool_param)[0], x, dout)

out, cache = max_pool_forward_naive(x, pool_param)
dx = max_pool_backward_naive(dout, cache)

# Your error should be around 1e-12
print 'Testing max_pool_backward_naive function:'
print 'dx error: ', rel_error(dx, dx_num)
```

```
Testing max_pool_backward_naive function:
dx error: 3.27563758008e-12
```

## Fast layers

Making convolution and pooling layers fast can be challenging. To spare you the pain, we've provided fast implementations of the forward and backward passes for convolution and pooling layers in the file `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:

```
python setup.py build_ext --inplace
```

The API for the fast versions of the convolution and pooling layers is exactly the same as the naive versions that you implemented above: the forward pass receives data, weights, and parameters and produces outputs and a cache object; the backward pass receives upstream derivatives and the cache object and produces gradients with respect to the data and weights.

**NOTE:** The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the following:

```
In [9]: from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
        from time import time

        x = np.random.randn(100, 3, 31, 31)
        w = np.random.randn(25, 3, 3, 3)
        b = np.random.randn(25,)
        dout = np.random.randn(100, 25, 16, 16)
        conv_param = {'stride': 2, 'pad': 1}

        t0 = time()
        out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
        t1 = time()
        out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
        t2 = time()

        print 'Testing conv_forward_fast:'
        print 'Naive: %fs' % (t1 - t0)
        print 'Fast: %fs' % (t2 - t1)
        print 'Speedup: %fx' % ((t1 - t0) / (t2 - t1))
        print 'Difference: ', rel_error(out_naive, out_fast)

        t0 = time()
        dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
        t1 = time()
        dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
        t2 = time()

        print '\nTesting conv_backward_fast:'
        print 'Naive: %fs' % (t1 - t0)
        print 'Fast: %fs' % (t2 - t1)
        print 'Speedup: %fx' % ((t1 - t0) / (t2 - t1))
        print 'dx difference: ', rel_error(dx_naive, dx_fast)
        print 'dw difference: ', rel_error(dw_naive, dw_fast)
        print 'db difference: ', rel_error(db_naive, db_fast)
```

```
Testing conv_forward_fast:
Naive: 8.386237s
Fast: 0.027078s
Speedup: 309.707659x
Difference: 5.76441667407e-11
```

```
Testing conv_backward_fast:
Naive: 9.075142s
Fast: 0.023224s
Speedup: 390.763728x
dx difference: 1.37522088985e-11
dw difference: 4.97269198917e-13
db difference: 8.05806780804e-14
```

```
In [10]: from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print 'Testing pool_forward_fast:'
print 'Naive: %fs' % (t1 - t0)
print 'fast: %fs' % (t2 - t1)
print 'speedup: %fx' % ((t1 - t0) / (t2 - t1))
print 'difference: ', rel_error(out_naive, out_fast)

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print '\nTesting pool_backward_fast:'
print 'Naive: %fs' % (t1 - t0)
print 'speedup: %fx' % ((t1 - t0) / (t2 - t1))
print 'dx difference: ', rel_error(dx_naive, dx_fast)
```

Testing pool\_forward\_fast:

Naive: 0.218179s

fast: 0.004214s

speedup: 51.774201x

difference: 0.0

Testing pool\_backward\_fast:

Naive: 1.068016s

speedup: 40.932619x

dx difference: 0.0

## Sandwich layers

There are a couple common layer "sandwiches" that frequently appear in ConvNets. For example convolutional layers are frequently followed by ReLU and pooling, and affine layers are frequently followed by ReLU. To make it more convenient to use these common patterns, we have defined several convenience layers in the file `cs231n/layer_utils.py`. Lets grad-check them to make sure that they work correctly:

In [11]: **from cs231n.layer\_utils import conv\_relu\_pool\_forward, conv\_relu\_pool\_backward**

```
x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w,
b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w,
b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w,
b, conv_param, pool_param)[0], b, dout)

print 'Testing conv_relu_pool_forward:'
print 'dx error: ', rel_error(dx_num, dx)
print 'dw error: ', rel_error(dw_num, dw)
print 'db error: ', rel_error(db_num, db)
```

```
Testing conv_relu_pool_forward:
dx error: 7.39177599082e-09
dw error: 5.02782991993e-10
db error: 3.50315221012e-11
```

In [12]: **from cs231n.layer\_utils import conv\_relu\_forward, conv\_relu\_backward**

```
x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b, co
nv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b, co
nv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b, co
nv_param)[0], b, dout)

print 'Testing conv_relu_forward:'
print 'dx error: ', rel_error(dx_num, dx)
print 'dw error: ', rel_error(dw_num, dw)
print 'db error: ', rel_error(db_num, db)
```

```
Testing conv_relu_forward:
dx error: 2.26033376732e-09
dw error: 2.02817906612e-10
db error: 2.39239728675e-11
```

```
In [13]: from cs231n.layer_utils import affine_relu_forward, affine_relu_backward

x = np.random.randn(2, 3, 4)
w = np.random.randn(12, 10)
b = np.random.randn(10)
dout = np.random.randn(2, 10)

out, cache = affine_relu_forward(x, w, b)
dx, dw, db = affine_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w, b)
[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w, b)
[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w, b)
[0], b, dout)

print 'Testing affine_relu_forward:'
print 'dx error: ', rel_error(dx_num, dx)
print 'dw error: ', rel_error(dw_num, dw)
print 'db error: ', rel_error(db_num, db)

Testing affine_relu_forward:
dx error:  4.21086180119e-10
dw error:  3.49094679052e-10
db error:  1.70177182826e-11
```

```
In [ ]:
```